

Endpoint	Feature	Model(s)	7d 2xx Requests	Monthly Volume	Prompt Token Method	Est. Tokens / Request	Est. Monthly Cost (USD)	P50 Latency (ms)	P95 Latency (ms)	Opportunity Flags	Ticket / Notes
/api/docs/summarize	Document Intelligence	claude-sonnet-4-5 (2 calls)	49	210.00	ceil(static chars + 4); ceil(avg response chars + 4)	5,058	\$11.61	7,468	8,490	Parallelize calls 1+2; est. ~3,062 ms p50 / ~3,481 ms p95. Anthropic cache not eligible: 857–858 static tokens < 1,024 minimum.	CLGO-14 — https://saurav-jvtq8dy.atlassian.net/browse/CLGO-14
/api/support/triage	Chat Assistant	claude-sonnet-4-5; claude-haiku-4-5 (2 calls)	56	240.00	ceil(static chars + 4); ceil(avg response chars + 4)	2,166	\$3.42	6,360	7,624	No parallelization: call 2 uses classification and call 3 uses draft. Cache not eligible (Sonnet 861 < 1,024; Haiku 36/31 < 4,096).	Slow, but no actionable opportunity; no ticket.
/api/chat/respond	Chat Assistant	gpt-4o	84	360.00	ceil(static chars + 4); ceil(avg response chars + 4)	1,570	\$3.20	1,382	1,617	Retry loop counts as one logical request. OpenAI caching is automatic; 907 static tokens are below its 1,024-token threshold.	Commented alternate gpt-4o call excluded.
/api/docs/translate	Document Intelligence	gemini-2.5-flash	28	120.00	ceil(static chars + 4); ceil(avg response chars + 4)	723	\$0.21	2,049	2,468	No second billable call; glossary fetch is non-LLM. Gemini caching is provider-managed, so no code-change flag.	No ticket.
/api/chat/summarize-thread	Chat Assistant	gpt-4o-mini (2 calls)	35	150.00	ceil(static chars + 4); ceil(avg response chars + 4)	742	\$0.06	2,470	2,950	Not parallelizable: title call consumes the rewrite output.	No ticket.
/api/onboarding/welcome-email	Growth	gpt-4o-mini	20	85.71	ceil(static chars + 4); ceil(avg response chars + 4)	343	\$0.02	1,118	1,284	None.	2 duplicate request rows excluded.
/api/onboarding/subject-line	Growth	gpt-4o-mini	0	0.00	ceil(static chars + 4); no response sample	12	\$0.00	0	0	No traffic in the window; volume, cost, and latency set to zero.	No ticket.

Feature	Endpoints	Monthly Volume	Est. Monthly Cost (USD)	Opportunities
Document Intelligence	2	330.00	\$11.82	1 real parallelization opportunity: /api/docs/summarize
Chat Assistant	3	750.00	\$6.68	No real opportunity: summarize-thread and triage chains are dependent
Growth	2	85.71	\$0.02	No opportunity
TOTAL	7	1,165.71	\$18.52	1 actionable opportunity; 1 Jira ticket

Provider	Model	Input Price	Output Price	Official Source	Checked Date	Notes
OpenAI	gpt-4o	\$2.50 / 1M text tokens	\$10.00 / 1M text tokens	https://developers.openai.com/api/docs/models/gpt-4o	Run-date basis 2026-07-16; re-verified 2026-08-04	Exact model alias; cached input \$1.25/1M; automatic caching.
OpenAI	gpt-4o-mini	\$0.15 / 1M text tokens	\$0.60 / 1M text tokens	https://developers.openai.com/api/docs/models/gpt-4o-mini	Run-date basis 2026-07-16; re-verified 2026-08-04	Exact model alias; cached input \$0.075/1M; automatic caching.
Anthropic	claude-sonnet-4-5	\$3.00 / 1M tokens	\$15.00 / 1M tokens	https://claude.com/pricing	Run-date basis 2026-07-16; re-verified 2026-08-04	Exact model; 5m cache write \$3.75/1M, read \$0.30/1M; 1,024-token minimum; cache_control required.
Anthropic	claude-haiku-4-5	\$1.00 / 1M tokens	\$5.00 / 1M tokens	https://claude.com/pricing	Run-date basis 2026-07-16; re-verified 2026-08-04	Exact model; 5m cache write \$1.25/1M, read \$0.10/1M; 4,096-token minimum; cache_control required.
Google	gemini-2.5-flash	\$0.30 / 1M text tokens	\$2.50 / 1M output tokens	https://ai.google.dev/gemini-api/docs/pricing	Run-date basis 2026-07-16; re-verified 2026-08-04	Exact model; standard paid tier; output includes thinking tokens.
OpenAI	text-embedding-3-small	\$0.02 / 1M input tokens	N/A	https://developers.openai.com/api/docs/models/text-embedding-3-small	Run-date basis 2026-07-16; re-verified 2026-08-04	Exact model; listed for completeness, excluded from generative attribution.
OpenAI	omni-moderation-latest	Free	Free	https://developers.openai.com/api/docs/models/omni-moderation-latest	Run-date basis 2026-07-16; re-verified 2026-08-04	Exact model alias; moderation models are free and excluded from attribution.

Run Date	Analysis Window	Commit SHA	Status / Exclusions
2026-07-16	2026-07-09 00:00:00 UTC – 2026-07-15 23:59:59 UTC	f4c8b6faed53bda1f3c75f6719a5dcbc731b9904	Complete (scheduled run date 2026-07-16; executed 2026-08-04). Estimated monthly generative spend: \$18.52. request_logs: 340 total; 5 excluded — 2 duplicate request_id rows (/api/onboarding/welcome-email), 1 null latency (/api/chat/respond), 1 zero latency (/api/docs/summarize), 1 negative latency (/api/support/triage); 335 survived. llm_call_logs: 534 total; 3 excluded — 2 repeated sequence-1 rows (/api/chat/summarize-thread), 1 provider/model mismatch (OpenAI gpt-4o on /api/docs/translate sequence 99); 531 survived. Code exclusions: commented alternate gpt-4o call in app/routers/chat.py; flag-off /api/onboarding/experimental-rewrite. Retry loop counted once per logical request. Dynamic request-time inputs and prior-call outputs excluded from static prompt-token estimates. Non-generative calls separated. Hardcoded test credential fallback detected at app/services/llm_clients.py:6; value not recorded. Jira: CLGO-14 created, unassigned, To Do — https://saurav-jvfqu8dy.atlassian.net/browse/CLGO-14 Teams: duplicate search found no Sheet/window match; update posted 2026-08-04 — https://teams.microsoft.com/l/message/19%3Aa8c987a2505040ca927477bbfb9d86eb%40thread.tacv2/1785829214278?groupId=3518d9d6-5c77-42fa-9459-8f9d17045638&tenantId=ad8b1651-7782-4309-a773-8b4fddfeab9f&createdTime=1785829214278&parentMessageId=1785829214278 Security note: Supabase reports RLS disabled on public.request_logs and public.llm_call_logs; no schema change made.

Endpoint	Call Seq	Provider	Model	Static Prompt	Static Chars	Est. Static Prompt Tokens	Avg Response Chars	Est. Response Tokens	Est. Cost / Request (USD)	Est. Monthly Cost Contribution (USD)	Dependency / Cache Notes
/api/chat/respond		1 OpenAI	gpt-4o	BASE_SYSTEM_PROMPT + RESPOND_SUFFIX	3,628	907	2,650.00	663	\$0.008898	\$3.20	Retry loop = one logical request. OpenAI automatic caching; static estimate below 1,024 tokens.
/api/chat/summarize-thread		1 OpenAI	gpt-4o-mini	THREAD_REWRITE_PROMPT	153	39	1,543.53	386	\$0.000237	\$0.04	Produces rewrite consumed by call 2.
/api/chat/summarize-thread		2 OpenAI	gpt-4o-mini	THREAD_TITLE_PROMPT	87	22	1,178.86	295	\$0.000180	\$0.03	Depends on call 1 output; not parallelizable.
/api/docs/summarize		1 Anthropic	claude-sonnet-4-5	BASE_SYSTEM_PROMPT + DOC_SUMMARY_SUFFIX	3,429	858	7,096.73	1,775	\$0.029199	\$6.13	Independent of call 2; parallelize. Static prefix below 1,024 cache minimum.
/api/docs/summarize		2 Anthropic	claude-sonnet-4-5	BASE_SYSTEM_PROMPT + DOC_TAG_SUFFIX	3,428	857	6,270.00	1,568	\$0.026091	\$5.48	Independent of call 1; parallelize. Static prefix below 1,024 cache minimum.
/api/docs/translate		1 Google	gemini-2.5-flash	TRANSLATE_PROMPT	137	35	2,750.00	688	\$0.001731	\$0.21	Only billable call; glossary fetch is non-LLM. Gemini caching provider-managed.
/api/onboarding/subject-line		1 OpenAI	gpt-4o-mini	Literal system prompt	48	12	0.00	0	\$0.000002	\$0.00	No traffic; output estimate unavailable and set to zero.
/api/onboarding/welcome-email		1 OpenAI	gpt-4o-mini	WELCOME_EMAIL_PROMPT	126	32	1,242.50	311	\$0.000191	\$0.02	Single call; no opportunity.
/api/support/triage		1 Anthropic	claude-sonnet-4-5	BASE_SYSTEM_PROMPT + CLASSIFY_SUFFIX	3,441	861	2,156.79	540	\$0.010683	\$2.56	Classification; call 2 depends on it. Static prefix below 1,024 cache minimum.
/api/support/triage		2 Anthropic	claude-haiku-4-5	DRAFT_RESPONSE_PROMPT	142	36	1,500.00	375	\$0.001911	\$0.46	Depends on call 1 classification; 4,096-token Haiku cache minimum.
/api/support/triage		3 Anthropic	claude-haiku-4-5	TONE_CHECK_PROMPT	121	31	1,290.00	323	\$0.001646	\$0.40	Depends on call 2 draft; not parallelizable; 4,096-token Haiku cache minimum.

Endpoint	Feature	Call Type	Model	7d 2xx Requests	Monthly Volume	P50 Latency (ms)	P95 Latency (ms)	Notes
/api/search/semantic	Search	Embedding	text-embedding-3-small	35	150.00	598	704	Listed separately; embedding cost is outside this generative attribution.
/api/moderation/check	Trust & Safety	Moderation	omni-moderation-latest	28	120.00	775	913	Listed separately; moderation is free and not billed per token.