

Endpoint	Feature	Model(s)	7d 2xx Requests	Monthly Volume	Prompt Token Method	Est. Tokens / Request	Est. Monthly Cost (USD)	P50 Latency (ms)	P95 Latency (ms)	Opportunity Flags	Ticket / Notes
/api/docs/summarize	Document Intelligence	claude-sonnet-4-5	49	210	cell(chars / 3); static input + observed output	6742.24	\$15.4774	7468	8490.2	PARALLEL: independent calls; ~3082 ms/request; CACHE: uncached Anthropic static prefix; up to \$1.30/mv CLGO-15 https://saurav-jytku8dy.atlassian.net/browse/CLGO-15 Two independent Sonnet calls, Anthropic cad	
/api/support/triage	Chat Assistant	claude-sonnet-4-5, claude-haiku-4-l	56	240	cell(chars / 3); static input + observed output	2885.43	\$4.5519	6359.5	7624.25	CACHE: uncached Anthropic static prefix; up to \$0.74/month; DEPENDENT: sequential outputs; not parallel CLGO-1 https://saurav-jytku8dy.atlassian.net/browse/CLGO-1 Existing issue refreshed; caching only; depende	
/api/chat/respond	Chat Assistant	gpt-4o	84	360	cell(chars / 3); static input + observed output	2094	\$4.2714	1381.5	1617.1	OPENAI AUTO-CACHE: eligible; no change required	Retry loop = one logical call; OpenAI prefix caching is automatic
/api/docs/translate	Document Intelligence	gemini-2.5-flash	28	120	cell(chars / 3); static input + observed output	968	\$0.2789	2049	2468.45		Already concurrent with glossary lookup; only one LLM call
/api/chat/summarize-thread	Chat Assistant	gpt-4o-mini	35	150	cell(chars / 3); static input + observed output	988.46	\$0.0836	2470	2950.1	DEPENDENT: sequential outputs; not parallelizable	Call 2 uses call 1 output; not parallelizable
/api/onboarding/welcome-email	Growth	gpt-4o-mini	20	85.71	cell(chars / 3); static input + observed output	456.5	\$0.0219	1116	1283.95		No actionable optimization
/api/onboarding/subject-line	Growth	gpt-4o-mini	0	0	cell(chars / 3); static input + observed output	16	\$0.0000	0	0		Discovered in pinned code; zero requests during the reporting window
TOTAL (billable only)			272	1165.71			\$24.6830			3 opportunities across 2 endpoints	

Feature	Endpoints	Monthly Volume	Est. Monthly Cost (USD)	Opportunities
Document Intelligence	2	330	\$15.7543	2
Chat Assistant	3	750	\$8.9069	1
Growth	2	85.71	\$0.0219	0
TOTAL	7	1165.71	\$24.6830	3

Provider	Model	Input Price (USD / 1M)	Output Price (USD / 1M)	Official Source	Checked Date	Notes
openai	gpt-4o	\$2.50	\$10.00	https://developers.openai.com/api/docs/models/gpt-4o	2026-07-16 report; retrieved 2026-08-	Exact model; standard USD per 1M tokens. Automatic eligible prompt caching is provider-managed; hit rate is unobserved, so base input pr
openai	gpt-4o-mini	\$0.15	\$0.60	https://developers.openai.com/api/docs/models/gpt-4o-mini	2026-07-16 report; retrieved 2026-08-	Exact model; standard USD per 1M tokens.
anthropic	claude-sonnet-4-5	\$3.00	\$15.00	https://platform.claude.com/docs/en/about-claude/pricing	2026-07-16 report; retrieved 2026-08-	Exact provider-listed Claude Sonnet 4.5 equivalent for API identifier claude-sonnet-4-5; first-party standard pricing; cache read \$0.30/MTok.
anthropic	claude-haiku-4-5	\$1.00	\$5.00	https://platform.claude.com/docs/en/about-claude/pricing	2026-07-16 report; retrieved 2026-08-	Exact provider-listed Claude Haiku 4.5 equivalent for API identifier claude-haiku-4-5; first-party standard pricing.
google	gemini-2.5-flash	\$0.30	\$2.50	https://ai.google.dev/gemini-api/docs/pricing#gemini-2.5-flash	2026-07-16 report; retrieved 2026-08-	Exact model; standard text USD per 1M tokens; output price includes thinking tokens.

Run Date	Analysis Window	Commit SHA	Status / Exclusions
2026-07-16	2026-07-09 00:00:00 UTC – 2026-07-15 23:59:59 UTC	f4c8b6faed53bda1f3c75f6719a5dcbc731b9904	RUN: Historical replay performed 2026-08-04; report as-of 2026-07-16; projected generative spend \$24.68/month; 7 billable endpoints.
2026-07-16	2026-07-09 00:00:00 UTC – 2026-07-15 23:59:59 UTC	f4c8b6faed53bda1f3c75f6719a5dcbc731b9904	REQUEST_LOG_SCREEN: 340 rows read; 5 excluded (null latency=1; zero/negative latency=2; duplicate request_id rows=2); 335 screened successful request rows retained.
2026-07-16	2026-07-09 00:00:00 UTC – 2026-07-15 23:59:59 UTC	f4c8b6faed53bda1f3c75f6719a5dcbc731b9904	LLM_CALL_LOG_SCREEN: 534 rows read; 4 excluded (repeated call_sequence_number rows=2; provider/model not in endpoint code=1; call linked to rejected duplicate request=1); 530 retained.
2026-07-16	2026-07-09 00:00:00 UTC – 2026-07-15 23:59:59 UTC	f4c8b6faed53bda1f3c75f6719a5dcbc731b9904	TOKEN_METHOD: Conservative character-count upper-bound heuristic: estimate one token per three Unicode characters, rounding each prompt/output upward; static hardcoded prompt text.
2026-07-16	2026-07-09 00:00:00 UTC – 2026-07-15 23:59:59 UTC	f4c8b6faed53bda1f3c75f6719a5dcbc731b9904	CODE_EXCLUSION: app/routers/chat.py:34 commented-out OpenAI alternate call was excluded.
2026-07-16	2026-07-09 00:00:00 UTC – 2026-07-15 23:59:59 UTC	f4c8b6faed53bda1f3c75f6719a5dcbc731b9904	CODE_EXCLUSION: app/routers/onboarding.py:9 ENABLE_EXPERIMENTAL_REWRITE=False; /api/onboarding/experimental-rewrite generative branch excluded.
2026-07-16	2026-07-09 00:00:00 UTC – 2026-07-15 23:59:59 UTC	f4c8b6faed53bda1f3c75f6719a5dcbc731b9904	ZERO_VOLUME: /api/onboarding/subject-line is executable in the pinned snapshot and retained with zero volume, latency, and monthly cost.
2026-07-16	2026-07-09 00:00:00 UTC – 2026-07-15 23:59:59 UTC	f4c8b6faed53bda1f3c75f6719a5dcbc731b9904	NON_GENERATIVE: /api/search/semantic embeddings (35 calls) and /api/moderation/check moderation (28 calls) listed separately and excluded from spend.
2026-07-16	2026-07-09 00:00:00 UTC – 2026-07-15 23:59:59 UTC	f4c8b6faed53bda1f3c75f6719a5dcbc731b9904	DEPENDENCIES: /api/docs/summarize has two independent LLM calls (~3062 ms/request parallelization). /api/chat/summarize-thread and /api/support/triage consume prior call outputs and are not idempotent.
2026-07-16	2026-07-09 00:00:00 UTC – 2026-07-15 23:59:59 UTC	f4c8b6faed53bda1f3c75f6719a5dcbc731b9904	PROMPT_CACHING: Anthropic Sonnet static prompts estimate 1143-1147 tokens, above the provider's 1024-token threshold; cache_control is absent. OpenAI prompt caching is automatic at 128k tokens.
2026-07-16	2026-07-09 00:00:00 UTC – 2026-07-15 23:59:59 UTC	f4c8b6faed53bda1f3c75f6719a5dcbc731b9904	PRICING: Official provider pages retrieved 2026-08-04; report valuation as-of 2026-07-16; USD per 1M standard input/output tokens; no model substitutions required.
2026-07-16	2026-07-09 00:00:00 UTC – 2026-07-15 23:59:59 UTC	f4c8b6faed53bda1f3c75f6719a5dcbc731b9904	SECRETS_FOUND: Hardcoded test-only OpenAI credential-like fallback found at app/services/llm_clients.py:6; value intentionally omitted.
2026-07-16	2026-07-09 00:00:00 UTC – 2026-07-15 23:59:59 UTC	f4c8b6faed53bda1f3c75f6719a5dcbc731b9904	SECURITY_CRITICAL: Row-level security is disabled on public.request_logs and public.llm_call_logs; design and review policies before enabling RLS. No configuration was changed.
2026-07-16	2026-07-09 00:00:00 UTC – 2026-07-15 23:59:59 UTC	f4c8b6faed53bda1f3c75f6719a5dcbc731b9904	TEAMS_DUPLICATE_CHECK: Searched channel posts from 2026-07-09 through 2026-07-16; returned July 16 records have no readable body and Chrome confirmed the visible prior post was not a duplicate.
2026-07-16	2026-07-09 00:00:00 UTC – 2026-07-15 23:59:59 UTC	f4c8b6faed53bda1f3c75f6719a5dcbc731b9904	COMPLETE: 24.68 USD/month; 3 opportunities across 2 endpoints; https://saurav-jviqu8dy.atlassian.net/browse/CLGO-15 created unassigned in To Do; https://saurav-jviqu8dy.atlassian.net/browse/CLGO-16 created unassigned in To Do.

Endpoint	Feature	Call #	Provider	Model	Static Prompt Chars	Est. Prompt Tokens	Avg Response Chars	Est. Response Tokens	7d Screened Calls	Monthly Logical Calls	Input \$ / 1M	Output \$ / 1M	Monthly Cost (USD)	Cache Opportunity	Dependency / Notes
/api/docs/summarize	Document Intelligence	1	anthropic	claude-sonnet-4.5	3429	1143	7096.73	2365.96	49	210	3	15	\$8.1729	Anthropic cache_control absent; ~1143 static tokens; upper-bound \$0.65/mo	Summary and keyword-tag calls both use original document_text; neither consumes the other output.
/api/docs/summarize	Document Intelligence	2	anthropic	claude-sonnet-4.5	3429	1143	6270	2090.29	49	210	3	15	\$7.3545	Anthropic cache_control absent; ~1143 static tokens; upper-bound \$0.65/mo	Summary and keyword-tag calls both use original document_text; neither consumes the other output.
/api/chat/respond	Chat Assistant	1	openai	gpt-4o	3628	1210	2650	884	84	360	2.5	10	\$4.2714	Automatic OpenAI caching; no missing configuration	One logical generative call
/api/support/triage	Chat Assistant	1	anthropic	claude-sonnet-4.5	3441	1147	2156.79	718.93	96	240	3	15	\$3.4140	Anthropic cache_control absent; ~1147 static tokens; upper-bound \$0.74/mo	Draft consumes classification; tone check consumes draft. The three Anthropic calls form a dependency chain and are not parallelizable.
/api/support/triage	Chat Assistant	2	anthropic	claude-haiku-4.5	142	48	1500	500.25	56	240	1	5	\$0.0118	None	Draft consumes classification; tone check consumes draft. The three Anthropic calls form a dependency chain and are not parallelizable.
/api/support/triage	Chat Assistant	3	anthropic	claude-haiku-4.5	121	41	1290	430.25	66	240	1	5	\$0.0261	None	Draft consumes classification; tone check consumes draft. The three Anthropic calls form a dependency chain and are not parallelizable.
/api/docs/translate	Document Intelligence	1	google	gemini-2.5-flash	152	51	2760	917	28	120	0.3	2.5	\$0.2769	None	One LLM call is already gathered concurrently with a PostgreSQL glossary lookup; no LLM-call pair exists.
/api/chat/summarize-thread	Chat Assistant	1	openai	gpt-4o-mini	153	51	1543.53	515.18	34	190	0.15	0.6	\$0.0475	None	Second title/handoff call consumes the rewritten summary produced by call 1; not parallelizable.
/api/chat/summarize-thread	Chat Assistant	2	openai	gpt-4o-mini	97	29	1178.86	393.29	35	190	0.15	0.6	\$0.0560	None	Second title/handoff call consumes the rewritten summary produced by call 1; not parallelizable.
/api/onboarding/welcome-email	Growth	1	openai	gpt-4o-mini	126	42	1242.5	414.5	20	85.71	0.15	0.6	\$0.0219	None	One logical generative call
/api/onboarding/subject-line	Growth	1	openai	gpt-4o-mini	48	16	0	0	0	0	0.15	0.6	\$0.0000	None	One logical generative call
TOTAL									467				\$24.6830		Response estimates use ceil(response_char_length / 3)

Endpoint	Feature	Call Type	Model	7d 2xx Requests	Monthly Volume	P50 Latency (ms)	P95 Latency (ms)	Attribution Exclusion
/api/search/semantic	Search	Embeddings	text-embedding-3-small	35	150	598	704	Embedding endpoint is non-generative and outside this attribution.
/api/moderation/check	Trust & Safety	Moderation	omni-moderation-latest	28	120	775	912.65	OpenAI moderation endpoint is non-generative and not billed per token.
TOTAL non-generative				63	270			Excluded from generative spend and ticket decisions