

Endpoint	Feature	Model(s)	7d 2xx Requests	Monthly Volume	Prompt Token Method	Est. Tokens / Request	Est. Monthly Cost (USD)	P50 Latency (ms)	P95 Latency (ms)	Opportunity Flags	Ticket / Notes
/api/docs/summarize	Document Intelligence	claude-sonnet-4-5 ×2	49	210.0	ceil(chars/4); static input + mean output	5058	\$11.6109	7,468.0	8,490.2	PARALLEL: independent summary + tags; ~3.08s mean theoretical critical-path saving. Anthropic cache candidate (conditional eligibility/hit rate; upper bound ~\$0.97/mo).	CLGO-16 (new)
/api/support/triage	Chat Assistant	claude-sonnet-4-5; claude-haiku-4-5 ×2	56	240.0	ceil(chars/4); static input + mean output	2166	\$3.4176	6,359.5	7,624.3	Anthropic stable classification cache candidate (conditional; upper bound ~\$0.56/mo). DEPENDENT classification → draft → tone check; not parallel.	CLGO-1 (updated)
/api/chat/respond	Chat Assistant	gpt-4o	84	360.0	ceil(chars/4); static input + mean output	1570	\$3.2031	1,381.5	1,617.1	OpenAI automatic caching for eligible prefixes; ~907 static tokens by heuristic is below 1,024. Retry loop counted as one logical call.	No ticket; below thresholds
/api/docs/translate	Document Intelligence	gemini-2.5-flash	28	120.0	ceil(chars/4); static input + mean output	723	\$0.2077	2,049.0	2,468.5	One billable call; glossary fetch is already concurrent. Gemini implicit caching is automatic when eligible; static prompt ~35 tokens.	No ticket; below thresholds
/api/chat/summarize-thread	Chat Assistant	gpt-4o-mini ×2	35	150.0	ceil(chars/4); static input + mean output	742	\$0.0627	2,470.0	2,950.1	DEPENDENT: title call consumes rewritten output; not parallel. Short static prompts; no manual OpenAI cache miss.	No ticket; below thresholds
/api/onboarding/welcome-email	Growth	gpt-4o-mini	20	85.7	ceil(chars/4); static input + mean output	343	\$0.0164	1,118.0	1,284.0	No actionable opportunity.	No ticket; below thresholds
/api/onboarding/subject-line	Growth	gpt-4o-mini	0	0.0	ceil(chars/4); static input + mean output	12	\$0.0000	0.0	0.0	No requests in window; retained at zero volume, cost, and latency.	No ticket; zero traffic

Feature	Endpoints	Monthly Volume	Est. Monthly Cost (USD)	Opportunities
Document Intelligence	2	330.0	\$11.8186	1 independent parallel pair; 1 conditional Anthropic cache candidate
Chat Assistant	3	750.0	\$6.6834	1 conditional Anthropic cache candidate; dependent chains explicitly not parallel
Growth	2	85.7	\$0.0164	None
TOTAL	7	1,165.7	\$18.5183	3 findings across 2 endpoints (1 parallel, 2 conditional cache candidates)

Provider	Model	Input Price	Output Price	Official Source	Checked Date	Notes
OpenAI	gpt-4o	\$2.50 / 1M tokens	\$10.00 / 1M tokens	https://developers.openai.com/api/docs/models/gpt-4o	2026-08-04	Standard text rate; exact alias. Requested pricing basis: 2026-07-16 retrospective window.
OpenAI	gpt-4o-mini	\$0.15 / 1M tokens	\$0.60 / 1M tokens	https://developers.openai.com/api/docs/models/gpt-4o-mini	2026-08-04	Standard text rate; exact alias.
Anthropic	claude-sonnet-4-5	\$3.00 / 1M tokens	\$15.00 / 1M tokens	https://platform.claude.com/docs/en/about-claude/pricing	2026-08-04	Standard Messages rate; exact family alias. 5m cache write \$3.75/MTok; hit \$0.30/MTok.
Anthropic	claude-haiku-4-5	\$1.00 / 1M tokens	\$5.00 / 1M tokens	https://platform.claude.com/docs/en/about-claude/pricing	2026-08-04	Standard Messages rate; exact family alias. 5m cache write \$1.25/MTok; hit \$0.10/MTok.
Google	gemini-2.5-flash	\$0.30 / 1M tokens	\$2.50 / 1M tokens	https://ai.google.dev/gemini-api/docs/pricing	2026-08-04	Paid standard text rate; output includes thinking tokens; exact model.
OpenAI	text-embedding-3-small	\$0.02 / 1M input tokens	N/A	https://developers.openai.com/api/docs/models/text-embedding-3-small	2026-08-04	Listed for completeness; excluded from generative attribution.
OpenAI	omni-moderation-latest	Free moderation endpoint	N/A	https://developers.openai.com/api/docs/models/omni-moderation-latest	2026-08-04	Not billed per token; excluded from generative attribution.

Run Date	Analysis Window	Commit SHA	Status / Exclusions
2026-08-04 (retrospective)	2026-07-09 00:00:00–2026-07-15 23:59:59 UTC	f4c8b6faed53bda1f3c75f6719a5dbc731b9904	Complete retrospective July 9–15 attribution: \$18.5183 estimated generative monthly spend; 7 billable endpoint rows (including zero-traffic subject-line). request_logs: 5/340 set aside (null 1, zero 1, negative 1, duplicate-ID rows 2). llm_call_logs: 3/534 set aside (provider/model mismatch 1, repeated-sequence rows 2); one further surviving call tied to an excluded request omitted from math. Commented alt_reply and flag-off experimental-rewrite excluded; embeddings/moderation separate. CLGO-16 created; CLGO-1 Summary/Description refreshed. Teams message 1785834321205 sent 2026-08-04 after no indexed Sheet/window match; historical bodies were blank and Chrome remained at launcher, so duplicate review was limited. See Methodology & Exclusions for details.

Endpoint	Feature	Sequence	Provider	Model	API Class	Static Prompt / Source	Static Chars	Est. Static Input Tokens	Successful Call Samples	Mean Response Chars	Est. Output Tokens	Est. Unit Cost (USD)	Est. Monthly Call Cost (USD)	Mean Call Latency (ms)
/api/docs/summarize	Document Intelligence	1	anthropic	claude-sonnet-4-5	Messages	BASE_SYSTEM_PROMPT + DOC_SUMMARY_SUFFIX	3429	858	49	7,096.7	1775	\$0.029199	\$6.131790	3,603.9
/api/docs/summarize	Document Intelligence	2	anthropic	claude-sonnet-4-5	Messages	BASE_SYSTEM_PROMPT + DOC_TAG_SUFFIX	3428	857	49	6,270.0	1568	\$0.026091	\$5.479110	3,078.4
/api/support/triage	Chat Assistant	1	anthropic	claude-sonnet-4-5	Messages	BASE_SYSTEM_PROMPT + CLASSIFY_SUFFIX	3441	861	56	2,156.8	540	\$0.010683	\$2.563020	2,180.2
/api/support/triage	Chat Assistant	2	anthropic	claude-haiku-4-5	Messages	DRAFT_RESPONSE_PROMPT	142	36	56	1,500.0	375	\$0.001911	\$0.469640	1,603.2
/api/support/triage	Chat Assistant	3	anthropic	claude-haiku-4-5	Messages	TONE_CHECK_PROMPT	121	31	56	1,290.0	323	\$0.001646	\$0.395040	1,410.7
/api/chat/respond	Chat Assistant	1	openai	gpt-4o	Chat Completions / Responses wrapper	BASE_SYSTEM_PROMPT + RESPOND_SUFFIX	3628	907	84	2,650.0	663	\$0.008898	\$3.203100	995.0
/api/docs/translate	Document Intelligence	1	google	gemini-2.5-flash	generateContent wrapper	TRANSLATE_PROMPT	137	35	28	2,750.0	688	\$0.001731	\$0.207660	1,462.5
/api/chat/summarize-thread	Chat Assistant	1	openai	gpt-4o-mini	Chat Completions / Responses wrapper	THREAD_REWRITE_PROMPT	153	39	34	1,543.5	386	\$0.000237	\$0.035618	944.3
/api/chat/summarize-thread	Chat Assistant	2	openai	gpt-4o-mini	Chat Completions / Responses wrapper	THREAD_TITLE_PROMPT	87	22	35	1,178.9	295	\$0.000180	\$0.027045	716.7
/api/onboarding/welcome-email	Growth	1	openai	gpt-4o-mini	Chat Completions / Responses wrapper	WELCOME_EMAIL_PROMPT	126	32	20	1,242.5	311	\$0.000191	\$0.016406	754.8
/api/onboarding/subject-line	Growth	1	openai	gpt-4o-mini	Chat Completions / Responses wrapper	Literal concise subject-line instruction	48	12	0	0.0	0	\$0.000002	\$0.000000	0.0

Endpoint	Feature Tag	Provider	Model	API Class	Screened 7d 2xx	Projected Monthly Volume	Attribution Treatment
/api/search/semantic	Search	openai	text-embedding-3-small	Embeddings	35	150	Separately listed; embedding cost excluded from generative attribution. p50 598 ms; p95 704 ms.
/api/moderation/check	Trust & Safety	openai	omni-moderation-latest	Moderation	28	120	Separately listed; moderation is free/not token-billed. p50 775 ms; p95 912.65 ms.

Topic	Detail
Run scope	Retrospective execution 2026-08-04 for requested 2026-07-09 00:00:00 through 2026-07-15 23:59:59 UTC window. Every Supabase data query used this bound.
Pinned source	main latest commit at/before 2026-07-16 17:05 IST (11:35 UTC): f4c8b6faed53bda1f3c75f6719a5dcbc731b9904, committed 11:21:05 UTC. Later a5622e6 at 12:00 UTC excluded.
Feature mapping	app.include_router tags win over legacy prefix. In particular /api/support/triage is Chat Assistant, not a separate Support feature.
Request screen	340 window rows; 5 set aside: 1 null total_latency_ms, 1 zero, 1 negative, and both rows for 1 duplicated request_id. Union excluded 5; 335 survive, all 2xx. Duplicate IDs excluded in full.
Call screen	534 window rows; 3 set aside: 1 provider/model mismatch (docs/translate sequence 99 openai gpt-4o) and both rows in 1 repeated request_id/call_sequence_number group. Union excluded 3; 531 survive. One surviving call attached to an excluded duplicate request was also omitted from successful-request cost math, leaving 530 linked call rows.
Latency / volume	p50 and p95 use continuous linear-interpolated percentiles on screened 2xx request total_latency_ms. Monthly endpoint requests = screened 7-day 2xx count ×30/7. No-traffic billable route retained at zero.
Token estimate	Rough heuristic: ceil(character count / 4). Static hardcoded system/instruction text only on input; successful screened call mean response_char_length /4 rounded up on output. No request-time user/document/context or prior-call output input tokens are available, so this is a directional static-plus-output lower-bound, not invoice reconciliation.
Cost formula	Per logical call = (estimated static input tokens × standard input \$/MTok + estimated output tokens × standard output \$/MTok)/1,000,000. Multiply by endpoint projected monthly request volume. Retry loop is one logical call.
Parallelism	docs/summarize summary and tag calls both read document_text and neither consumes the other's output: real independent pair. Paired shorter-call mean 3,078 ms (p50 3,062; p95 3,481) is theoretical critical-path saving. chat/summarize-thread and support/triage chains consume prior outputs and are explicitly not parallel.
Caching	Anthropic helper has no explicit cache_control. Stable Sonnet prompts are candidates; rough 90%-read upper bounds ~\$0.97/month docs and ~\$0.56 /month support exclude cache writes, TTL and miss rates. Character/4 estimates 857–861 are below Anthropic Sonnet's 1,024-token minimum; validate exact tokenizer, viable stable prefix and temporal hit rate before claiming savings. Haiku short prompts are far below its 4,096 minimum. OpenAI and Gemini implicit caching is automatic for eligible prefixes; no equivalent manual-enable miss was counted.
Excluded code	Commented-out alt_reply in app/routers/chat.py is disabled; /api/onboarding/experimental-rewrite is gated by ENABLE_EXPERIMENTAL_REWRITE=False. Both excluded from billable inventory and cost.
Non-generative	Embeddings /api/search/semantic and moderation /api/moderation/check are in the separate Non-Generative Calls tab; neither contributes to the \$18.5183 generative total.
Secrets note	A hardcoded fake test-only credential fallback is present in app/services/llm_clients.py line 6. Value intentionally not reproduced. Repository README identifies this as a synthetic fixture, not production.
Security advisory	Supabase metadata reported RLS disabled on public.request_logs and public.llm_call_logs. No remediation was applied. Review exposure and design policies before enabling RLS; enabling without policies can block access.
Pricing basis	Official provider pages verified 2026-08-04 for the requested retrospective 2026-07-16 attribution. Standard first-party rates used; exact aliases were listed, no closest-equivalent substitution.
Ticket rule	Only endpoints with cost ≥\$25 or p95 ≥5,000 ms AND an opportunity: docs/summarize and support/triage qualify on p95. CLGO-16 new (closed CLGO-2 did not block); CLGO-1 updated Summary/Description only.
Diagnostic boundary	No repository code, provider configuration, caching setting, or database schema was changed.
Teams duplicate check	Listed channel posts from 2026-07-09 through actual send time 2026-08-04; exact Sheet-ID and July 9–15 searches returned no indexed match. History list/fetch returned blank bodies, and the Chrome deep link stayed at the Teams launcher, so the review could not inspect those bodies. No positive duplicate was found; message 1785834321205 was sent. This limitation is recorded rather than treating blank content as proof.